

ECE365: Data Structures and Algorithms II (DSA 2)

Course Intro and Review of DSA 1

The Basics

- Lectures will be in person, Wednesdays 4 pm – 6 pm in Rm. 505
- If we are forced to be remote at any point during the semester, those lectures will take place via Teams
- My webpage: <http://faculty.cooper.edu/sable2>
 - From there, you can find a link to the course webpage
 - I'll post syllabus info and assignments on the course webpage
- My email: carl.sable@cooper.edu

The Textbook

- “Data Structures and Algorithm Analysis in C++” by Mark Allen Weiss
 - I recommend the 4th Edition (updated for C++11)
 - This is the same book that was used for DSA 1
- Other references (not required):
 - “Algorithms”, by Sedgewick, the 3rd edition
 - “Introduction to Algorithms”, the 3rd edition, by Cormen, Leiserson, Rivest, and Stein

Grading

- The grading breakdown for the course is:
 - Four programs (50% total)
 - Two tests (50% total)
- All programs will be written in **C++**
- The first three programs will build off each other

What are Data Structures and Algorithms?

- Informal definitions:
 - **Data structures** are constructs for organizing data
 - **Algorithms** are methods of solving a problem
 - Or: An *algorithm* is a well-defined sequence of computational steps that transforms structured input into structured output
- We will see a more formal definition of an algorithm near the end of this course

Helpful Notations for Analyzing Algorithms

- **Big-O** (a.k.a. Big-Oh): $T(N) = O(f(N))$ if there exist positive constants c and n_0 such that $T(N) \leq c * f(N)$ when $N \geq n_0$
- **Big-Omega** (a.k.a. Omega): $T(N) = \Omega(g(N))$ if there exist positive constants c and n_0 such that $T(N) \geq c * g(N)$ when $N \geq n_0$
- **Big-Theta** (a.k.a. Theta): $T(N) = \Theta(h(N))$ if and only if $T(N) = O(h(N))$ and $T(N) = \Omega(h(N))$
- We also covered *Little-O*, but that is less commonly used
- Typically, we use these notations to analyze the *worst-case running time* (really, the *number of computational steps*) of an algorithm
- Sometimes we analyze the average-case running time or the memory requirements of an algorithm

Analyzing Pseudo-code or Code

- We covered a series of mathematical rules applying to Big-O or the related notations
- Example: if $T(N)$ is a polynomial of degree k , then $T(N) = \Theta(N^k)$
- We also went over a series of rules that tell us how to use these notations to analyze pseudo-code
- Example: work inside out and multiply when you are dealing with nested loops
- We also covered *the Master Theorem* which allows us to analyze certain recurrences that occur when we analyze some recursive algorithms
- We analyzed various solutions to the *Maximum Subsequence Sum Problem*
- We saw that algorithms with different Big-O running times take vastly different actual times to run
- We also explained why *exponential* running time algorithms are even much worse than high-order polynomial time algorithms

Overview of Overview of C++ (not a typo)

- DSA 1 included an *overview of C++*
- Concepts covered included streams, strings, pointers, dynamic memory, references, exception handling, overloaded functions, etc.
- We discussed **object-oriented programming** in C++
- Concepts covered include *classes, objects, encapsulation, information hiding, operator overloading, inheritance, polymorphism*, etc.
- We also covered C++ *templates*, which allows *generic programming*

Lists

- We discussed the notion of an **abstract data type** (ADT); this defines the values that can be stored and the types of operations that can be applied to the values
- An abstract data type does not imply how the actual data type is to be implemented
- We discussed the **list** abstract data type, and discussed advantages and disadvantages of implementing an *array-based list* vs. a *linked list*
- Arrays allow a constant time findKth operation, but require linear time for insertion and deletion since existing elements need to be shifted
- Linked lists allow constant time insertions and deletions (given access to the previous node), but findKth becomes a linear operation
- C++ provides a *vector* class that supports resizable arrays, and a *list* class that provides doubly linked list functionality (both implementations rely on templates)

Stacks and Queues

- We covered **stacks** and **queues**, two simple but very important ADTs
- I tend to think of a stack as a closed bucket and a queue as an open pipe
- You can insert into either data structure using a **push** or delete using a **pop**; most modern sources prefer the terms **enqueue** and **dequeue** when dealing with a queue
- A stack uses a **LIFO** strategy, and a queue uses a **FIFO** strategy
- Both stacks and queues can be implemented with arrays or linked lists
- One of your programming assignments involved linked lists, stacks, and queues; the implementation required the use of classes, inheritance, polymorphism, and templates
- We covered at least a couple of linear algorithms that use stacks; e.g., an algorithm for checking if an expression containing left- and right-markers is balanced
- We also discussed how every program you write with function calls relies on a **call stack** (a.k.a. **the stack**) of *activation records*; without this, *recursion* would not be possible

Trees

- We covered was **trees**, starting with terminology; e.g., root, leaf, internal node, child, parent, ancestor, descendant, depth, etc.
- Three traversal strategies, generally implemented with recursion, are *preorder traversal*, *postorder traversal*, and *inorder traversal*
- Another traversal strategy, generally implemented using a queue, is *breadth-first search*
- In a **binary tree**, every node can have at most two children

Binary Search Trees

- A **binary search** tree is a binary tree that constrains the ordering of nodes based on their *keys*
 - Binary search trees support average-case logarithmic time insertion and search
 - Items can be traversed in sorted order (according to keys) in linear time using an inorder traversal
 - An ordinary binary search tree has worst-case linear insertion and search
- An **AVL tree** is one type of **balanced binary search** tree that guarantees worst-case logarithmic time insertions and searches
 - For each node, v , of an AVL tree, T , the heights of the two subtrees of v must differ by at most 1
 - We saw how this property can be maintained without violating logarithmic time insertion by using single rotations and double rotations
- Other types of balanced trees include *splay trees*, *red black trees*, and *B+ trees* (not binary)

Sorting and Simple Sorts

- We spent a few weeks covering **sorting** algorithms
- Some of the sorts we covered are **stable**, meaning that items with equal keys remain in the same relative order before and after the sort
- Three simple, quadratic-time sorts are *bubble sort*, *selection sort*, and *insertion sort*
- Bubble sort and insertion sort can be very fast for sequences that start off in close-to-sorted order

$\Theta(N \log N)$ Sorting Algorithms

- Two $\Theta(N \log N)$ sorts that we covered in detail are **mergesort** and **quicksort**
- Both sorts typically rely on recursion and use a divide-and-conquer strategy
- Mergesort sorts the left half of the sequence, then sorts the right half of the sequence, then *merges* the two sorted halves together
- Quicksort relies on a *partition*, which uses a *pivot*; a common strategy for choosing a pivot is called *median-of-three partitioning*
- Quicksort is one of the fastest general sorting algorithms in practice
- Most implementations of quicksort are $\Theta(N^2)$ in the worst case, but this is exponentially unlikely with a good implementation
- Quicksort is not a stable sort, whereas mergesort is stable

Linear Sorting Algorithms

- We proved that $O(N \log N)$ is the best you can do for any **comparison-based sort**
- *Bin sort* (a.k.a. *bucket sort* or *counting sort*) is a linear time sort, but it makes very strict assumptions about the possible data
- (Least-significant-digit) **radix sort** is an extension of this type of sort that makes less strict assumptions about the data
- For example, radix sort it can be used to sort 32-bit integers and still provides a linear time sort

Indirect Sorting

- If you are sorting very large structures, or items in a linked list, it might be helpful to use **indirect sorting**
- Two methods of indirect sorting are a *pointer sort* or an *index sort* (the latter cannot be applied to link lists)
- If an *in-place sort* is necessary, you can follow an indirect sort with another routine to move items into their correct locations

Hash Tables and Hash Functions

- The last topic we covered was **hash tables**
- Hash tables typically support two or three operations: *insert*, *search*, and (possibly) *delete*
- The goal of a hash table is to provide the supported operations in *constant average time*
- A hash table typically makes use of a large array; it is often advisable to make the size of the array a prime number
- A hash table implementation must rely on a **hash function** which maps a **key** (often a string) to a slot in the hash table
- A *modular hash function*, as its last step, takes the remainder when some calculated number is divided by the size of the hash table

Collisions and Collision Resolution Strategies

- A **collision** occurs when non-equal keys get mapped by the hash function to the same location
- When this happens, the hash table must rely on a **collision resolution strategy**
 - The simplest collision resolution strategy is **separate chaining**, in which each slot stores a linked list of items that hash to that slot
 - For separate chaining, the *load factor* represents the average size of each list
 - Two other commonly used collision resolution strategies are **linear probing** and **double hashing**; these are examples of *open addressing schemes*
 - For these strategies, the load factor is the fraction of the table that is occupied, and it must be less than 1 (typically you want to keep it much lower; e.g., less than 0.5)
 - Open addressing schemes are more complicated and only support *lazy deletion*; however, they allow faster operations by avoiding dynamic memory allocation

Rehashing

- It is typically a good idea to keep the load factor small (e.g., less than 0.5), especially for the open addressing schemes
- One way to ensure this is with **rehashing**, which resizes the hash table when it fills to a certain point
 - Everything from the original table gets reinserted into the new, larger table
 - When you rehash, you eliminate the lazily deleted items
- Although this will mean that occasionally there will be a relatively slow (linear) insertion, they will still be fast (constant time) on average
- Even if you know the maximum amount of possible data, if there will usually be much less data, rehashing can avoid wasting memory

Applications of Hash Tables

- *Compilers* use hash tables to keep track of symbol tables
- Some *game-playing programs* use them for transposition tables
- *Database management systems* use them (or B+ trees, or related data structures) as indexes
- *Information retrieval systems* (e.g., web search) use them to map words to document lists
- *Spell checkers* use them to detect errors and highlight them